Build Generative AI Agents with Vertex AI Agent Builder

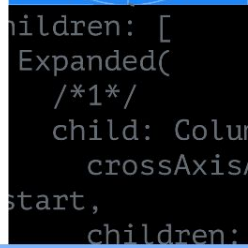

Google Developer Groups
Cloud • Rajkot

Cloud Community Day 2024

Introduction

- AI solutions that optimize business operations and improve decision-making processes.
- Expert in Google Cloud Platform technologies, including Vertex AI with a strong background in building and integrating intelligent conversation AI solutions for enhanced customer engagement
- M.Sc Physics from NIT Surat
- Linkedin Top AI Voice

Yash Kavaiya
AI/ ML Engineer at TCS

Agenda

- Introduction to Generative AI
- What are Large Language Models ?
- Neural Networks
- What are Large Language Models ?
- Overview of Vertex AI
- Retrieval-Augmented Generation
- Demo-1 Vertex AI Agent with Agent Builder
- Chainlit UI with Google Gemini API
- PDF Using Streamlit App and Gemini API

Google Developer Groups
Cloud • Rajkot

Cloud Community Day 2024

Generative AI

Create new content

(Audio, Code ,Text, Video, Images)

Automatically using computer program

Generative AI is a subset of Deep Learning

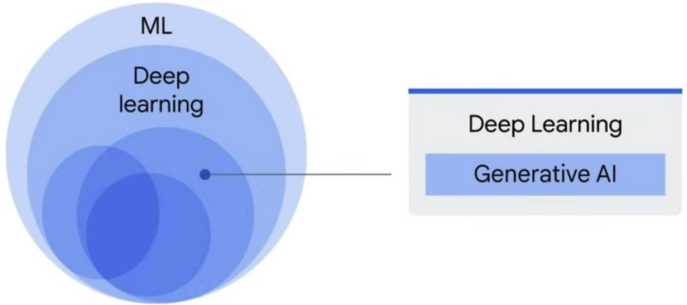

Predictive ML vs Generative AI

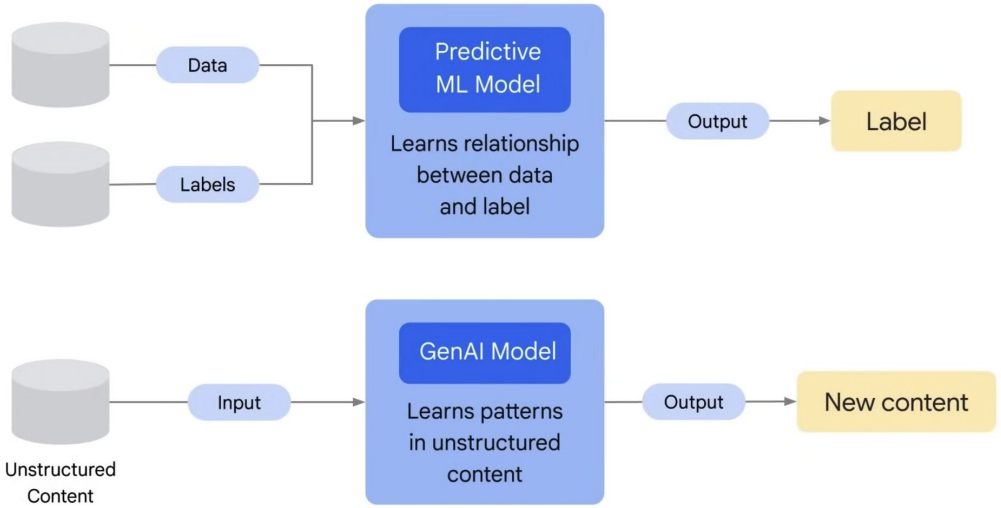

Discriminative AI vs Generative AI

Deep Learning Model Types

Discriminative

- Used to classify or predict
- Typically trained on a dataset of labeled data
- Learns the relationship between the features of the data points and the labels

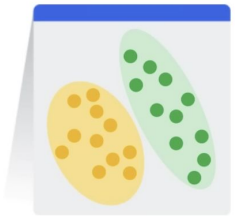

Generative

- Generates new data that is similar to data it was trained on
- Understands distribution of data and how likely a given example is
- Predict next word in a sequence

Gen AI Output

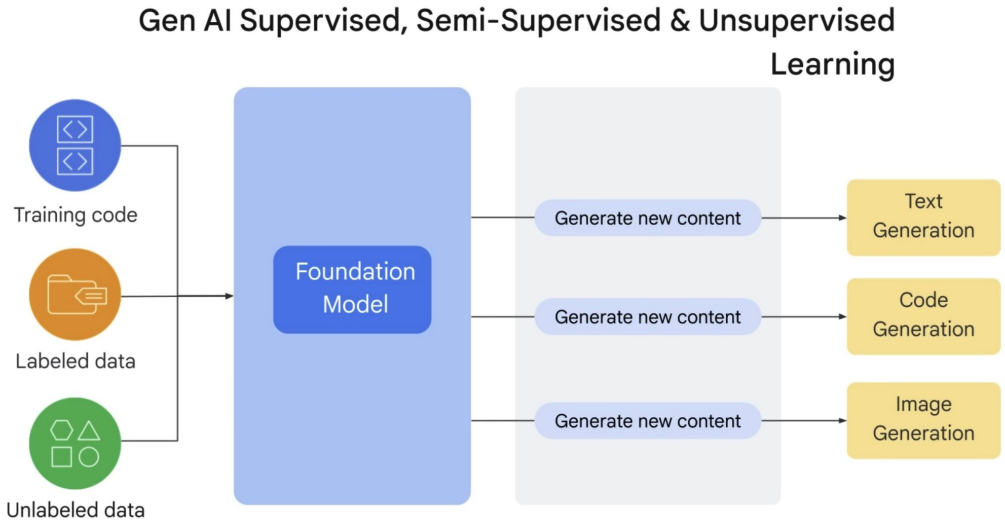

Neural Networks

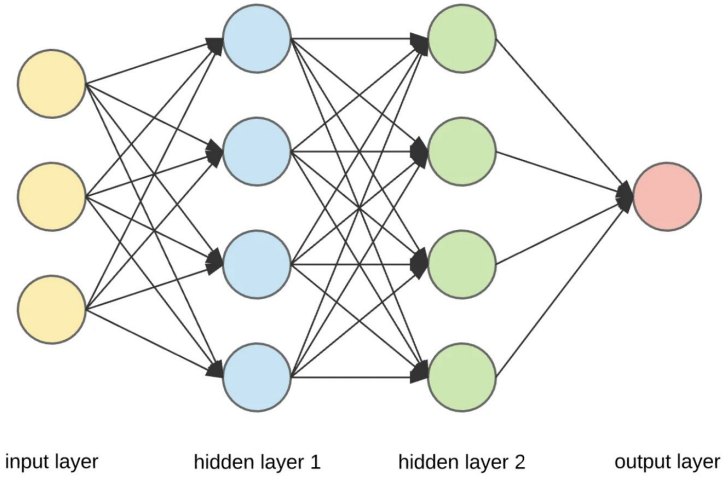

A neural network is a computational model inspired by the way biological neural networks in the human brain process information. It consists of interconnected nodes (neurons) organized in layers, which work together to recognize patterns and make decisions.

Neural Networks

Computer Vision

Convolutional NNs (+ResNets)

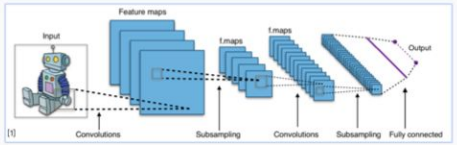

Natural Lang. Proc.

Recurrent NNs (+LSTMs)

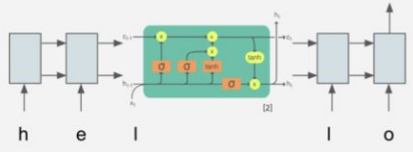

Speech

Deep Belief Nets (+non-DL)

Translation

Seq2Seq

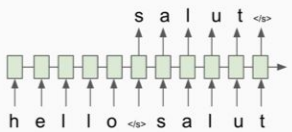

RL

BC/GAIL

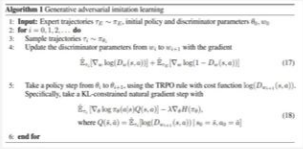

[1] CNN image CC-BY-SA by Aphex34 for Wikipedia https://commons.wikimedia.org/wiki/File:Typical_cnn.png
[2] RNN image CC-BY-SA by GChe for Wikipedia https://commons.wikimedia.org/wiki/File:The_LSTM_Cell.svg

Google Developer Groups
Cloud • Rajkot

Cloud Community Day 2024

Now : It’s all Transformers

Transformer cartoon (DALL-E)

Attention is all you need

- Positional Encoding
- Attention
- Self-attention

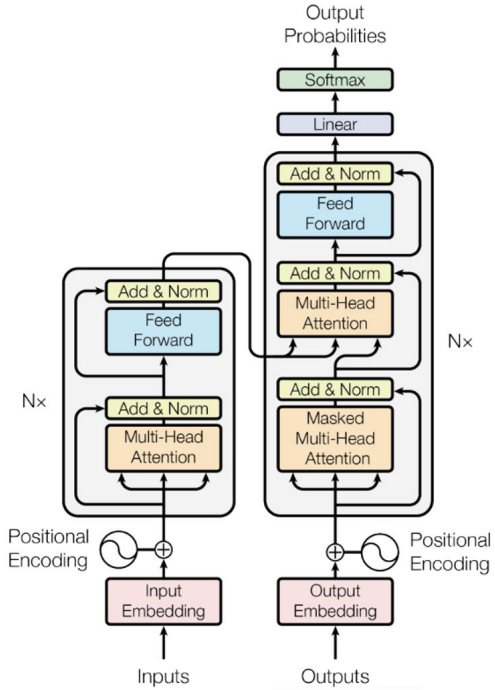

What are Large Language Models ?

- Type of machine learning models that are trained on a large data
- Generates outputs for tasks, such as text generation, question answering, and machine translation
- Based on deep learning neural networks, such as the Transformer Architecture
- At the core LLM’s are just language models that can predict the next word in a sentence

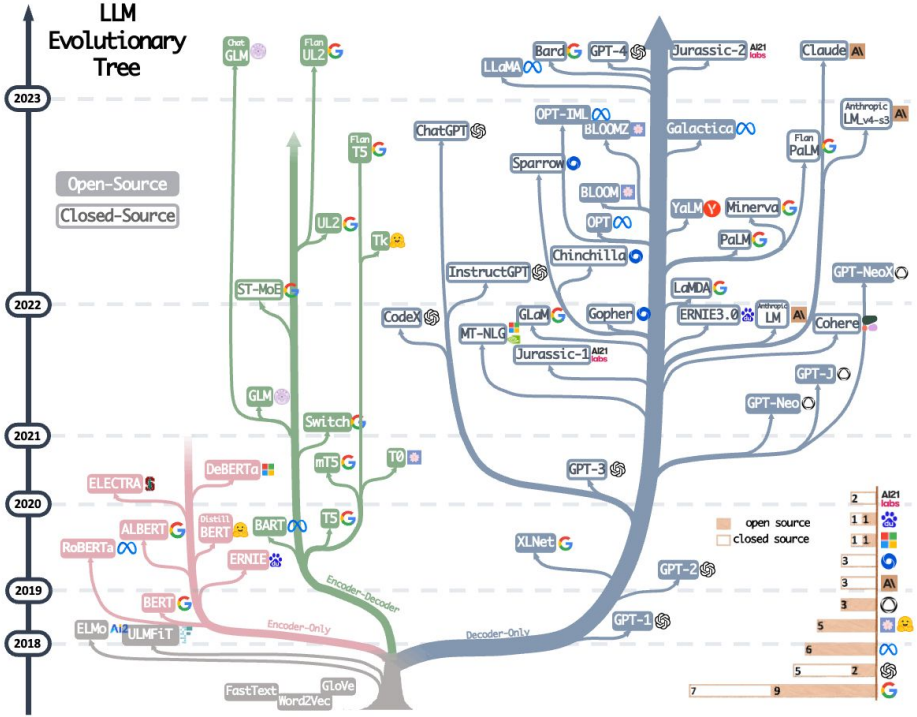

Google Developer Groups
Cloud • Rajkot

Cloud Community Day 2024

What is Vertex AI

Vertex AI is a Google Cloud platform that simplifies building and using artificial intelligence. It offers tools for every step, from data prep to deployment, and caters to both beginners and experts. Vertex AI includes pre-built models for easy use and also allows you to build your own.

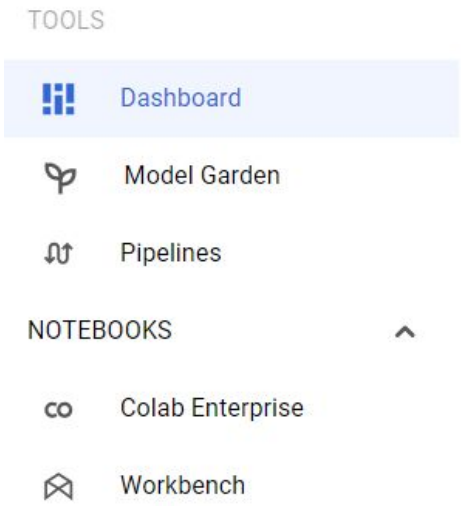

Model Garden

Vertex AI Model Garden is essentially a repository of pre-trained machine learning models that you can easily discover, use, and customize.

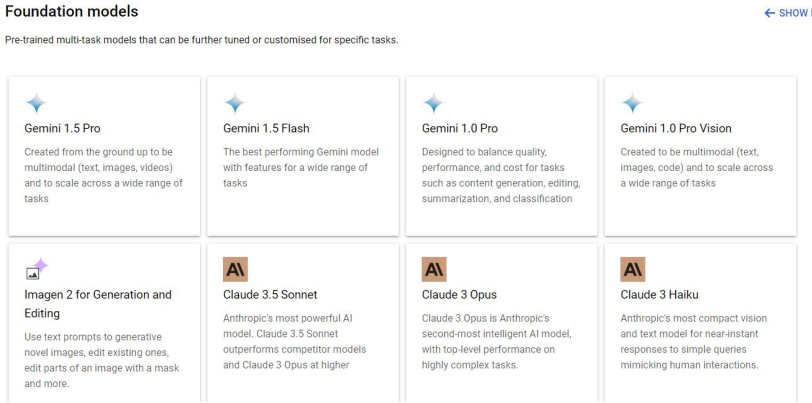

Google Developer Groups
Cloud • Rajkot

Cloud Community Day 2024

What's Inside the Garden?

- **Diverse Model Collection:** You'll find a vast array of models covering various domains like computer vision, natural language processing, and more.
- **Foundation Models:** These are powerful models that can be adapted to various tasks with minimal fine-tuning.
- **Task-Specific Models:** These models are tailored for specific tasks, like image classification or sentiment analysis.
- **APIs:** You can directly access and use model functionalities through APIs, without needing to delve into the underlying code.

Modalities	
Language	59
Vision	88
Tabular	7
Document	6
Speech	1
Video	4
Multimodal	17

How Does It Benefit You?

- Accelerated Development
- Experimentation
- MLOps Integration
- Customization

Tasks	
Generation	68
Classification	59
Detection	39
Extraction	24
Recognition	22
Translation	20
Embedding	7
Segmentation	8
Retrieval	2
Open vocabulary detection	2
Open vocabulary segmentation	2
Tracking	1
Forecasting	5

Tokens

A token is the smallest unit of text that a language model can process. It could be a word, a punctuation mark, or even a subword unit. For example, the sentence "The quick brown fox jumps over the lazy dog" would be broken down into the following tokens: "The", "quick", "brown", "fox", "jumps", "over", "the", "lazy", "dog".

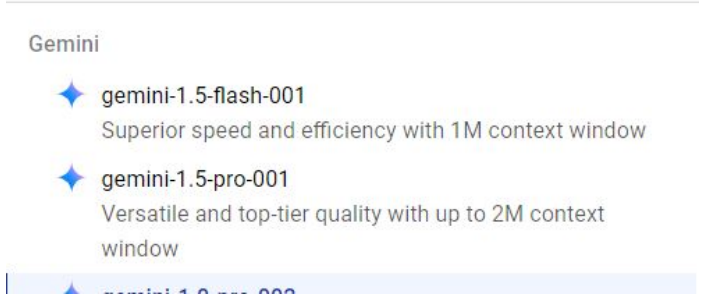

Embeddings

- An embedding is a numerical representation of a token.
- It's a vector (a list of numbers) that captures the semantic and syntactic meaning of the token.
- Words with similar meanings tend to have similar embeddings.

Word Embedding: "What is AI " - [0.42569, 0.82569, 0.385236, ...]

Embeddings

- Let’s consider an example with word embeddings.
- Suppose we have the words "king," "queen," "man," and "woman."
- In the embedding space, these words might have the following vectors:
  "king" == [0.5, 1.2, -0.3]
  "queen" == [0.6, 1.1, -0.2]
  "man" == [0.4, 0.8, -0.5]
  "woman" == [0.5, 0.9, -0.4]

These vectors allow us to perform arithmetic operations to uncover relationships, such as:

"king" − "man" + "woman" ≈ "queen"

Vector DB

A vector database is a specialized database designed to efficiently store, manage, and retrieve high-dimensional vector data. Unlike traditional databases that primarily deal with structured tabular data, vector databases excel at handling data points represented as vectors in a multi-dimensional space.

Vector DB

Key Functionalities

- Vector Storage: Efficiently stores high-dimensional vectors.
- Indexing: Creates indexes to optimize search performance.
- Similarity Search: Quickly finds vectors similar to a given query vector.
- Scalability: Handles large volumes of data and high query rates.

Milvus zilliz Pinecone

vespa marqo qdrant

Weaviate Chroma Vald

LanceDB

https://www.linkedin.com/pulse/choosing-vector-database-your-gen-ai-stack-abhinav-srivastava/

Temperature

Controls the randomness of the output.

Higher temperature: Increases randomness, leading to more diverse and creative outputs but potentially less coherent text.
Lower temperature: Decreases randomness, resulting in more focused and deterministic outputs but potentially less creative.

Top_k

Limits the number of possible next tokens to the top k most probable ones at each generation step.

Higher top_k: Increases the diversity of output by considering more options.
Lower top_k: Reduces the diversity but can improve focus and coherence.

Top_p

Definition: Selects the next token from the set of tokens whose cumulative probability exceeds a certain threshold (top_p).
Higher top_p: Increases the diversity of output by considering a larger set of tokens.
Lower top_p: Reduces diversity but can improve focus and coherence.

What is Vertex AI Agents

Vertex AI Agent Builder is a suite of tools from Google Cloud that simplifies building and deploying Gen AI agents. It caters to developers of all experience levels, offering

- Easily build no code conversational AI agents
- Ground in Google search and/or your enterprise data with our RAG offerings
- Rapidly create low-code to high-code AI application

Introducing LangChain

- An Open Source modular framework for building applications powered by language models
- Chatbots and virtual assistants
- Text generation and summarization
- Document question answering
- Relies on language models to reason
- Connects a language model to sources of context (prompts, contextual content)
- Combines components (language models, agents,
- memory stores) into complex workflows

Source:https://www.langchain.com/

Why we need Retrieval-Augmented Generation (RAG)

- Enhanced accuracy
- Contextual responses
- Reduced hallucinations
- Dynamic knowledge updates
- Cost efficiency

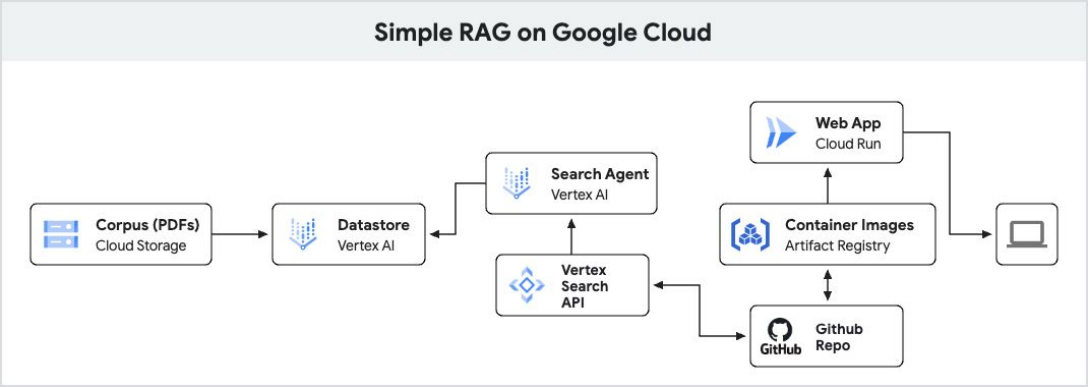

Google Developer Groups

Cloud Community Day 2024

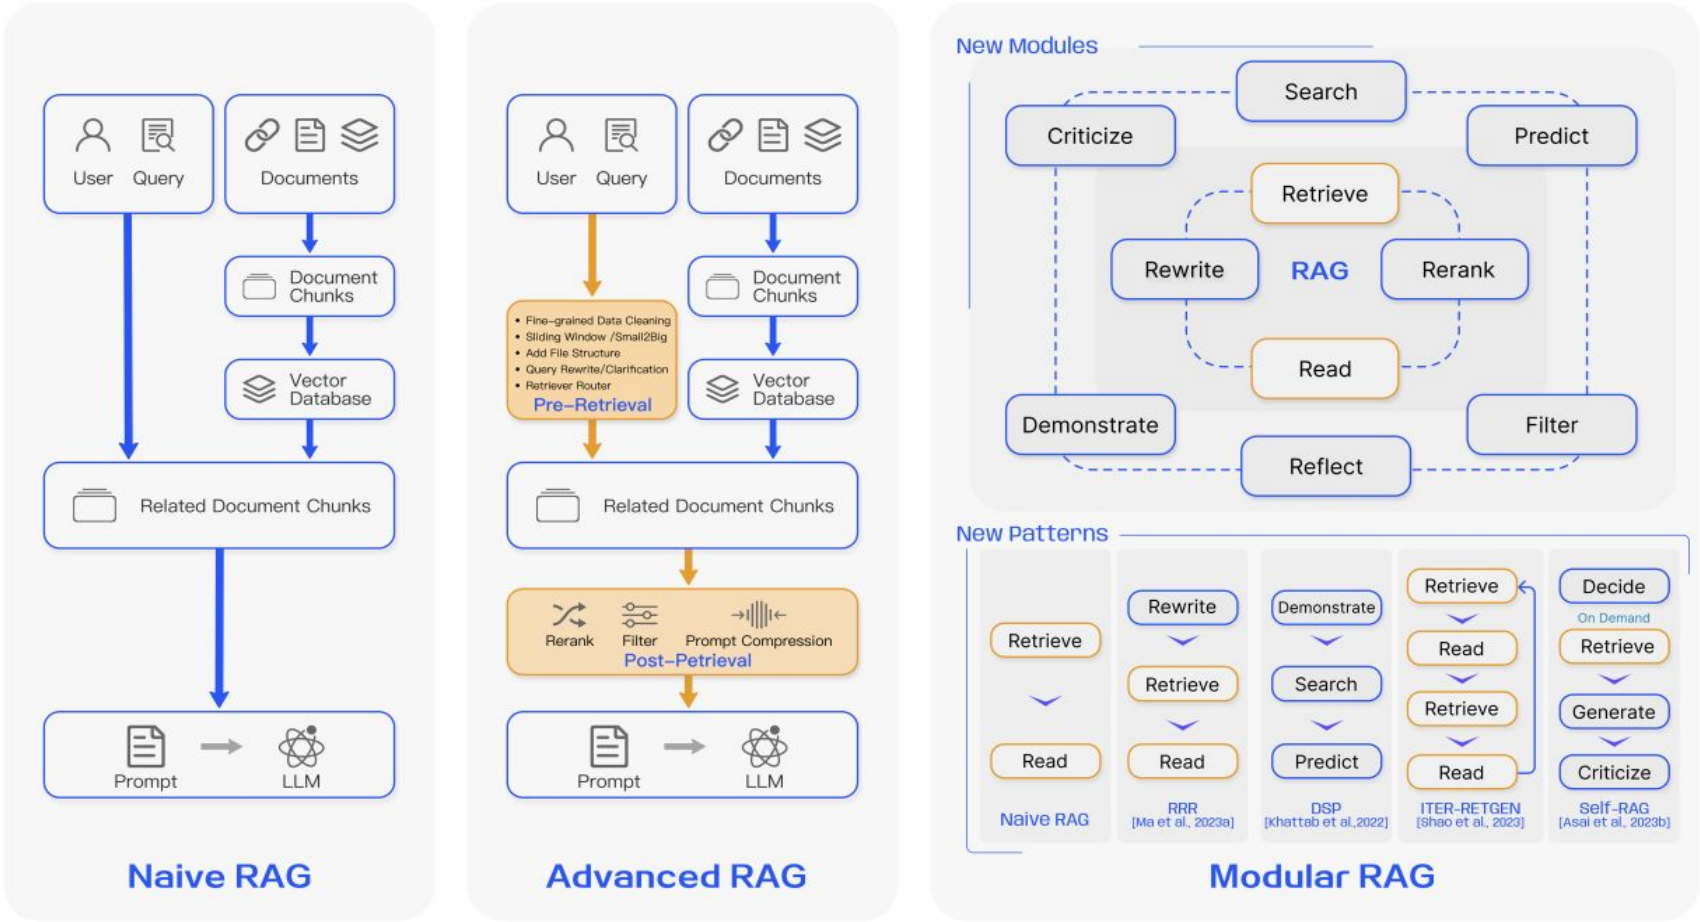
User
Query
Documents
Document Chunks
Vector Database
Related Document Chunks
Prompt
LLM
Naive RAG
• Fine-grained Data Cleaning
• Sliding Window /Small2Big
• Add File Structure
• Query Rewrite/Clarification
• Retriever Router
Pre-Retrieval
Rerank
Filter
Prompt Compression
Post-Petrieval
Advanced RAG
New Modules
Search
Criticize
Predict
Retrieve
Rewrite
RAG
Rerank
Read
Demonstrate
Filter
Reflect
New Patterns
Naive RAG
RRR
[Ma et al., 2023a]
DSP
[Khattab et al.,2022]
ITER-RETGEN
[Shao et al., 2023]
Decide
On Demand
Generate
Criticize
Self-RAG
[Asai et al., 2023b]
Modular RAG

Feature	Vertex AI Agent with RAG-based Chatbot	Vector Search Solution	Vertex AI Embedding with Vector DB
Overview	No code -low code tool	Implements a search solution using vector-based retrieval	Combines Vertex AI Embedding with a vector database to enhance search and recommendations
Components Used	Vertex AI Agent, Dialogflow, Cloud Run	Vertex AI Vector Search , Cloud Run	Vertex AI Embedding, Cloud Run / Compute Engine
Scalability	High, handles large datasets and multiple interactions concurrently	High, optimized for large-scale search operations	High, supports large-scale embedding and search operations
Customization	Low, customizable conversation flows and response logic	Moderate, focuses on search and retrieval customization	High, flexible embedding and search parameter tuning
Cost Efficiency	Moderate, depends on usage and integration complexity	High, cost-effective for large-scale search operations	High, optimized for cost-efficient embedding and search solutions

Popular Python Framework for Quick Development

Streamlit

gradio

Chainlit

mesop

Google Developer Groups
Cloud • Rajkot

Cloud Community Day 2024

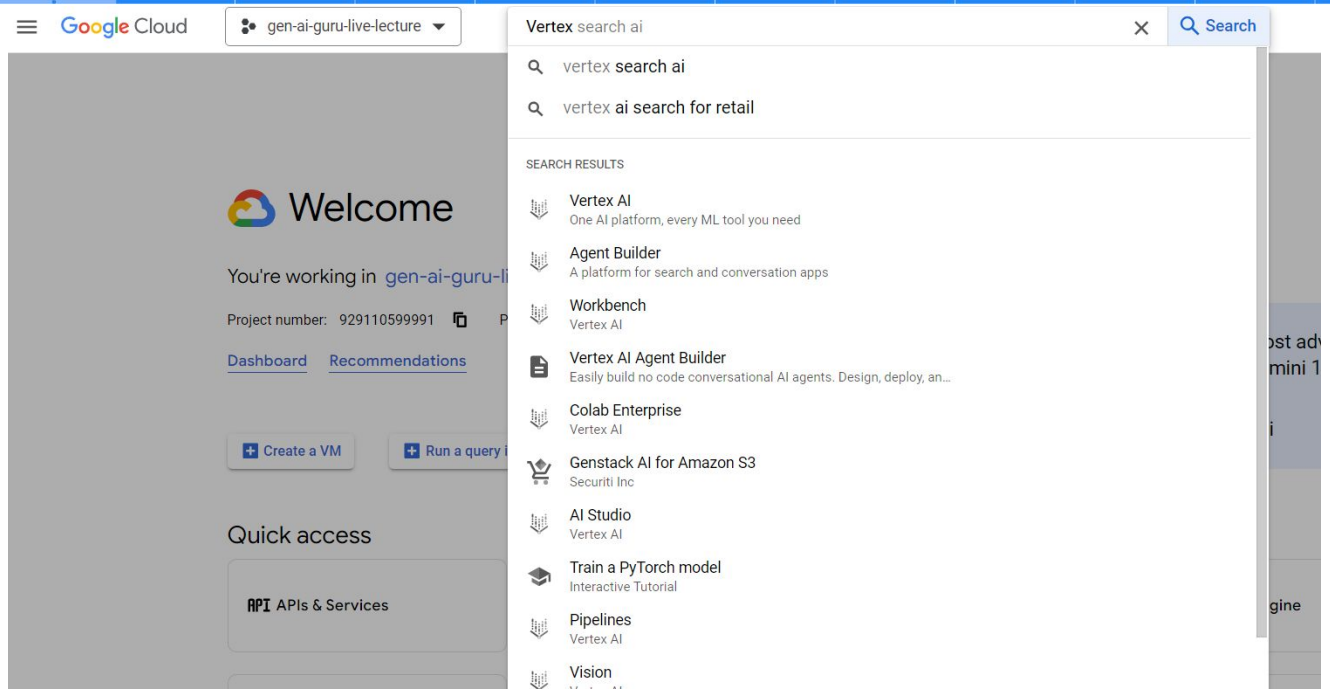

Google Developer Groups
Cloud • Rajkot

Cloud Community Day 2024

Enabling Vertex AI API

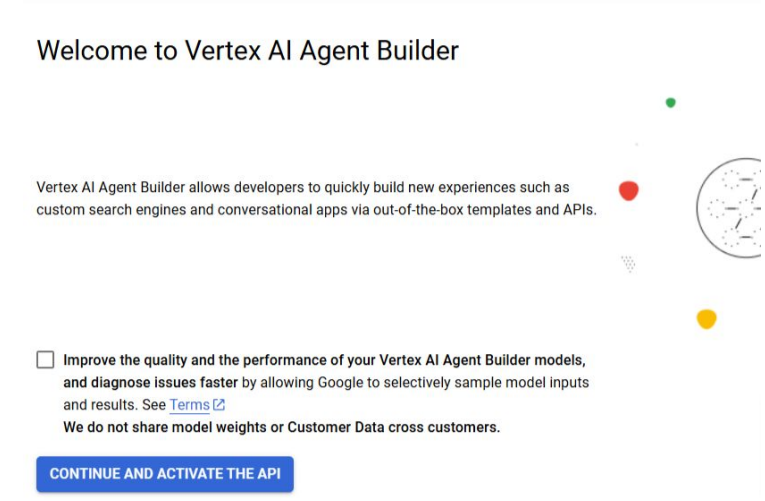

Google Developer Groups
Cloud • Rajkot

Cloud Community Day 2024

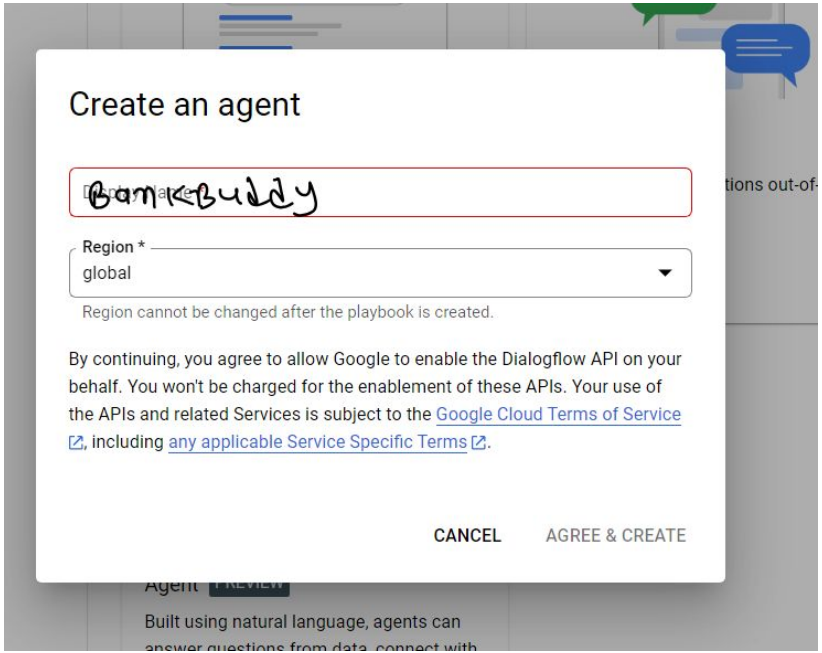

Google Developer Groups

Cloud Community Day 2024

Google Cloud | gen-ai-guru-live-lecture | vertex ai search and conversation | Search | 12

Agent Builder

Create App

LEARN

- Apps
- Data Stores
- Monitoring
- Settings

1. Type
2. Configuration
3. Data

Select app type

Select the type of application you want to create

Search

Get quality results out-of-box and easily customize the engine

SELECT

Chat

Answer complex questions out-of-the-box

SELECT

Recommendations

Create a content recommendation engine

SELECT

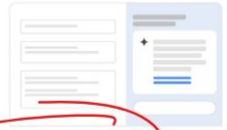

Agent PREVIEW

Built using natural language, agents can answer questions from data, connect with

Agent Console

App: BankBuddy

BankBuddy | Version history | Save

Basics | Examples

Agent name*: BankBuddy

An agent is the basic building block of a Vertex AI Conversation app. Each agent is defined to handle specific tasks. Learn more

Goal

Goal*: Help customers to Solve their banking, investment and FAQ questions.

High level description of the goal the agent intends to accomplish. Learn more

Instructions

Sample

Instructions

```
- Greet the user, then ask how you can help them today.
- Summarize the user's request and ask them to confirm that you understood
correctly.
- If necessary, seek clarifying details.
- Use ${TOOL: Example tool name} to help the user with their task.
- Use ${AGENT: Example agent name} to help the user with a complex subtask.
- Thank the user for their business and say goodbye.
```

Preview agent: BankBuddy

Send a message to see how your agent responds

Teach your agent by saving examples with intended responses Learn more

Agent: BankBuddy

Select generative model: gemini-1.0-pro-001

Note this is the model you are testing with. To change the model used by your published agent, visit Settings

Enter user input

Live Demo

Code

Google Developer Groups
Cloud • Rajkot

Cloud Community Day 2024

Demo -1

Vertex AI Agent with Agent Builder

Google Developer Groups
Cloud • Rajkot

Cloud Community
Day 2024

Demo -2

Chainlit UI with Google Gemini API

Google Developer Groups
Cloud • Rajkot

Cloud Community
Day 2024

Demo -3

PDF Using Streamlit App and Gemini API

Google Developer Groups
Cloud • Rajkot

Cloud Community
Day 2024

Google Developer Groups

Cloud • Rajkot

Cloud Community Day 2024

Thanks

https://linktr.ee/yashkavaiya